

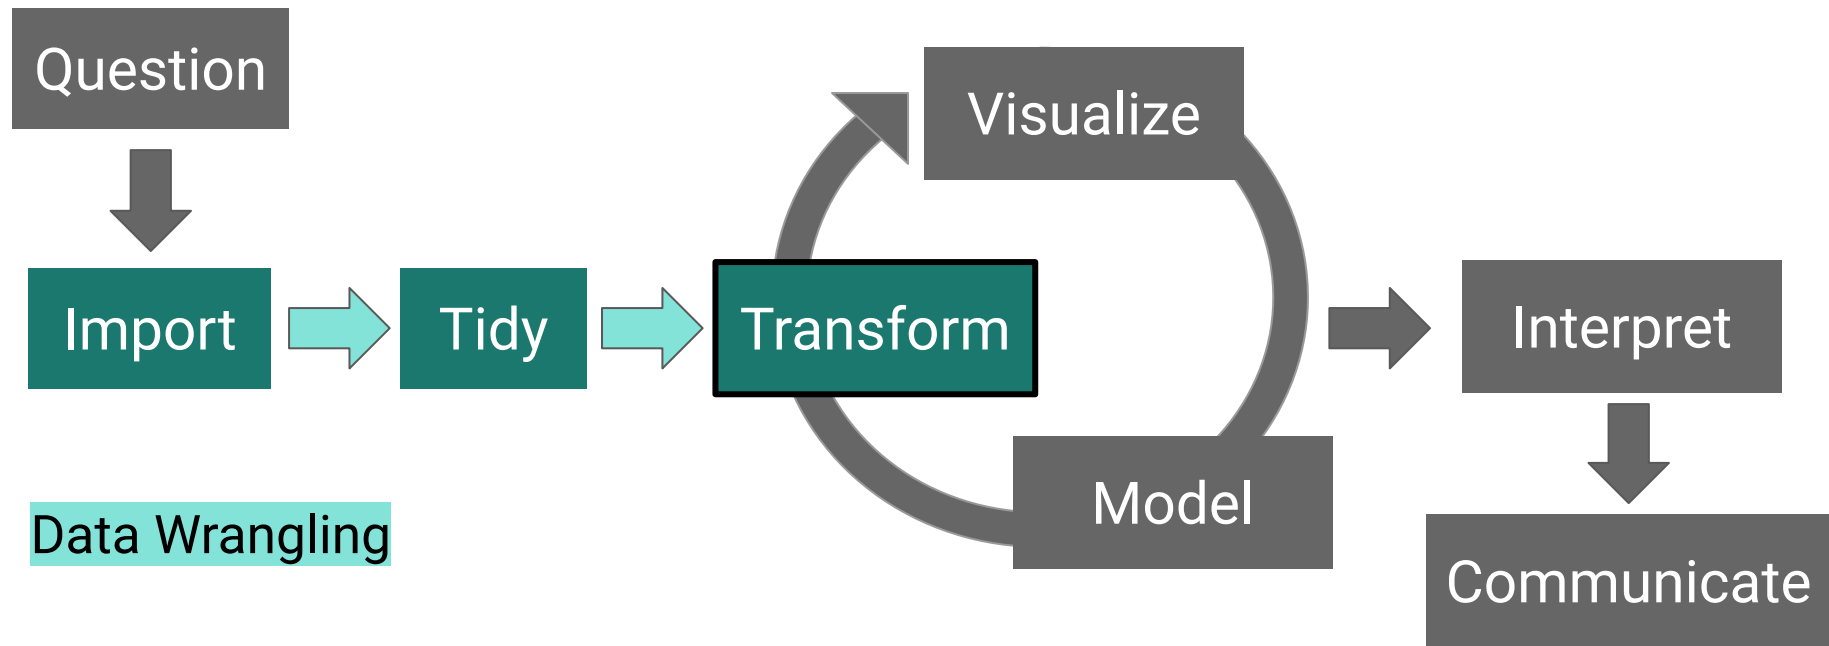
Merging Data

Lecture 6

Objective

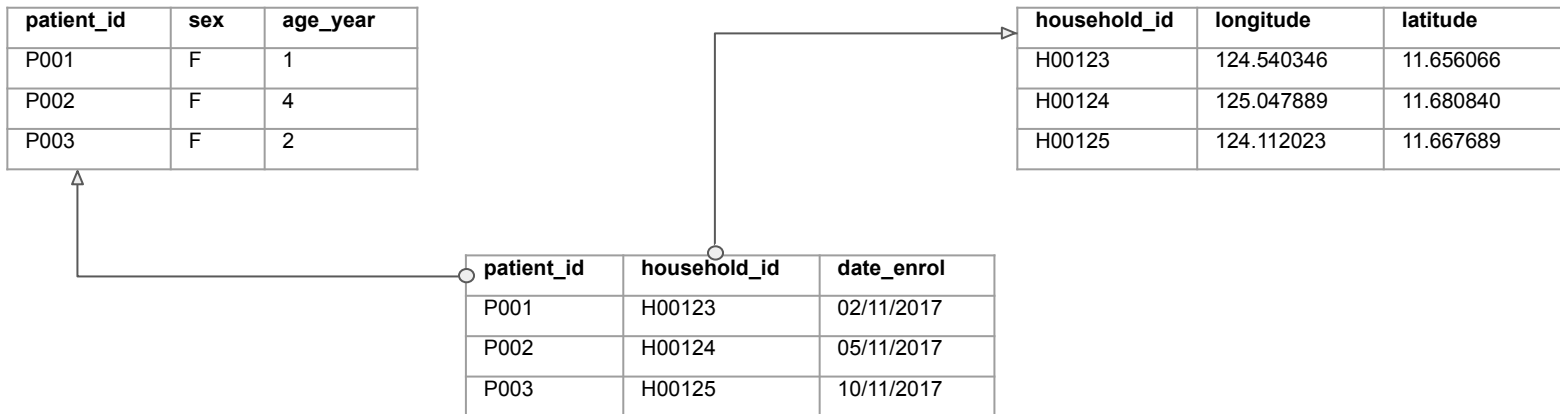
- To merge data frames

Motivation




Motivation

- it is rare that data analysis involves only a single table of data
- usually many tables of data are involved, which must combine to answer the questions of interest
- most data stored in relational database (e.g. SQL database)



Keys

- variables used to connect each pair of tables
- variables that uniquely identifies an observation



3 Functions for Merging Data Frames

1. **Mutating joins** - add new variables to one data frame that match observations in another
2. **Filtering joins** - filter observations from one data frame that match an observation in the other table
3. **Set operations** - treat observations as elements of a set


Mutating joins

- mutating join allows you to combine variables from two tables
- first matches observations by their keys, then copies across variables from one table to the other
- inner join and outer join

inner_join()

- inner join matches pairs of observations whenever their keys are equal
- keeps observations that appear in both tables

```
> inner_join(x, y, by = c("A", "B"))
```



inner_join()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104

```
> inner_join(data_1, data_2, by = c("patient_ID" = "patient_id"))
```

```
# A tibble: 4 x 5
```

patient_ID	sex	age_year	weight_kg	height_cm
<chr>	<chr>	<dbl>	<dbl>	<dbl>
1 P001	female	1	9.1	73
2 P002	female	4	16.4	NA
3 P003	female	2	10.5	85
4 P004	male	3	13.2	95

left_join()

- outer join keeps observations that appear in at least one of the tables
- **left join** keeps all observations in x

```
> left_join(x, y, by = c("A", "B"))
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	D
a	t	3
b	u	2
d	w	1



A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA

left_join()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104

```
> left_join(data_1, data_2, by = c("patient_ID" = "patient_id"))
```

```
# A tibble: 5 x 5
```

patient_ID	sex	age_year	weight_kg	height_cm
<chr>	<chr>	<dbl>	<dbl>	<dbl>
1 P001	female	1	9.1	73
2 P002	female	4	16.4	NA
3 P003	female	2	10.5	85
4 P004	male	3	13.2	95
5 P005	male	4	NA	NA

right_join()

- outer join keeps observations that appear in at least one of the tables
- **right join** keeps all observations in y

```
> right_join(x, y, by = c("A", "B"))
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	D
a	t	3
b	u	2
d	w	1



A	B	C	D
a	t	1	3
b	u	2	2
d	w	NA	1

right_join()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104

```
> right_join(data_1, data_2, by = c("patient_ID" = "patient_id"))
```

```
# A tibble: 5 x 5
```

patient_ID	sex	age_year	weight_kg	height_cm
<chr>	<chr>	<dbl>	<dbl>	<dbl>
1 P001	female	1	9.1	73
2 P002	female	4	16.4	NA
3 P003	female	2	10.5	85
4 P004	male	3	13.2	95
5 P006	NA	NA	15.9	104

full_join()

- **full join** keeps all observations in x and y

```
> full_join(x, y, by = c("A", "B"))
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	D
a	t	3
b	u	2
d	w	1



A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA
d	w	NA	1

full_join()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104


```
> full_join(data_1, data_2, by = c("patient_ID" = "patient_id"))
```

patient_ID	sex	age_year	weight_kg	height_cm
<chr>	<chr>	<dbl>	<dbl>	<dbl>
1 P001	female	1	9.1	73
2 P002	female	4	16.4	NA
3 P003	female	2	10.5	85
4 P004	male	3	13.2	95
5 P005	male	4	NA	NA
6 P006	NA	NA	15.9	104

Filtering joins

- filtering joins match observations in the same way as mutating joins, but affect the observations, instead of the variables
- **semi_join()** - **keeps** all observations in x that have a match in y

```
> semi_join(x, y)
```



semi_join()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104

```
> semi_join(data_1, data_2, by = c("patient_ID" = "patient_id"))
```

```
# A tibble: 4 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3

anti_join()

- **anti_join()** - **drops** all observations in x that have a match in y
- useful to see observations that are not joined

```
> anti_join(x, y)
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	D
a	t	3
b	u	2
d	w	1



A	B	C
c	v	3

anti_join()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104

```
> anti_join(data_1, data_2, by = c("patient_ID" = "patient_id"))
```

```
# A tibble: 1 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P005	male	4

Set operations

- use for combining rows
- these operations work with a complete row, comparing the values of every variable
- expect the x and y inputs to have the same variables, and treat the observations like sets
- **intersect()** - return only observations in both x and y
- **union()** - return unique observations in x and y
- **setdiff()** - return observations in x, but not in y

intersect()

- returns rows that appear only in both x and y

```
> intersect(x, y)
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	C
c	v	3
d	w	4



A	B	C
c	v	3

intersect()

```
> data_x
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96
3	P003	female	2	10.5	85

```
> data_y
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P003	female	2	10.5	85
2	P004	male	3	13.2	95
3	P005	male	4	15.9	104

```
> intersect(data_x, data_y)
```

```
# A tibble: 1 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P003	female	2	10.5	85

union()

- returns rows that appear in x or y
- duplicates will be removed

> union(x, y)


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	C
c	v	3
d	w	4



A	B	C
a	t	1
b	u	2
c	v	3
d	w	4

union()

```
> data_x
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96
3	P003	female	2	10.5	85

```
> data_y
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P003	female	2	10.5	85
2	P004	male	3	13.2	95
3	P005	male	4	15.9	104

```
> union(data_x, data_y)
```

```
# A tibble: 5 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96
3	P003	female	2	10.5	85
4	P004	male	3	13.2	95
5	P005	male	4	15.9	104

setdiff()

- returns rows that appear in x but not in y

```
> setdiff(x, y)
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	C
c	v	3
d	w	4



A	B	C
a	t	1
b	u	2

setdiff()

```
> data_x
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96
3	P003	female	2	10.5	85

```
> data_y
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P003	female	2	10.5	85
2	P004	male	3	13.2	95
3	P005	male	4	15.9	104

```
> setdiff(data_x, data_y)
```

```
# A tibble: 2 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96

bind_rows()

- to stack tables on top of the other as they are

```
> bind_rows(x, y)
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	C
c	v	3
d	w	4



A	B	C
a	t	1
b	u	2
c	v	3
c	v	3
d	w	4

bind_rows()

```
> data_x
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96
3	P003	female	2	10.5	85

```
> data_y
```

```
# A tibble: 3 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P003	female	2	10.5	85
2	P004	male	3	13.2	95
3	P005	male	4	15.9	104

```
> bind_rows(data_x, data_y)
```

```
# A tibble: 6 x 5
```

	patient_ID	sex	age_year	weight_kg	height_cm
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	P001	female	1	9.1	73
2	P002	female	4	16.4	96
3	P003	female	2	10.5	85
4	P003	female	2	10.5	85
5	P004	male	3	13.2	95
6	P005	male	4	15.9	104

bind_cols()

- to paste tables beside each other as they are

```
> bind_cols(x, y)
```


x

A	B	C
a	t	1
b	u	2
c	v	3



y

A	B	D
a	t	3
b	u	2
d	w	1



A	B	C	A	B	D
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

bind_cols()

```
> data_1
```

```
# A tibble: 5 x 3
```

patient_ID	sex	age_year
<chr>	<chr>	<dbl>
1 P001	female	1
2 P002	female	4
3 P003	female	2
4 P004	male	3
5 P005	male	4

```
> data_2
```

```
# A tibble: 5 x 3
```

patient_id	weight_kg	height_cm
<chr>	<dbl>	<dbl>
1 P001	9.1	73
2 P002	16.4	NA
3 P003	10.5	85
4 P004	13.2	95
5 P006	15.9	104

```
> bind_cols(data_1, data_2)
```

```
# A tibble: 5 x 6
```

patient_ID	sex	age_year	patient_id	weight_kg	height_cm
<chr>	<chr>	<dbl>	<chr>	<dbl>	<dbl>
1 P001	female	1	P001	9.1	73
2 P002	female	4	P002	16.4	NA
3 P003	female	2	P003	10.5	85
4 P004	male	3	P004	13.2	95
5 P005	male	4	P006	15.9	104

Take-away message

- Knowing how to merge tables is useful since data are not always available in a single table